


## Gradient Boosting


- Supervised learning
  - Finds the best function  $f$  to minimize the cost
- Contains one base value (constant) and to this it adds weak learners to better the prediction
  - The constant value is  $F_0$  (our initial best guess)
- Basically just adding a new trained model to our best guess to get a better guess

$$\hat{F}(x) = \sum_{m=1}^M \gamma_m h_m(x) + \text{const}$$

$$F_0 + \nu \gamma_1 = F_1$$



$$F_1 + \nu \gamma_2 = F_2$$



<https://www.youtube.com/watch?v=lOwsMpdjxog>

---

## XGBoost (Extreme Gradient Boosting)

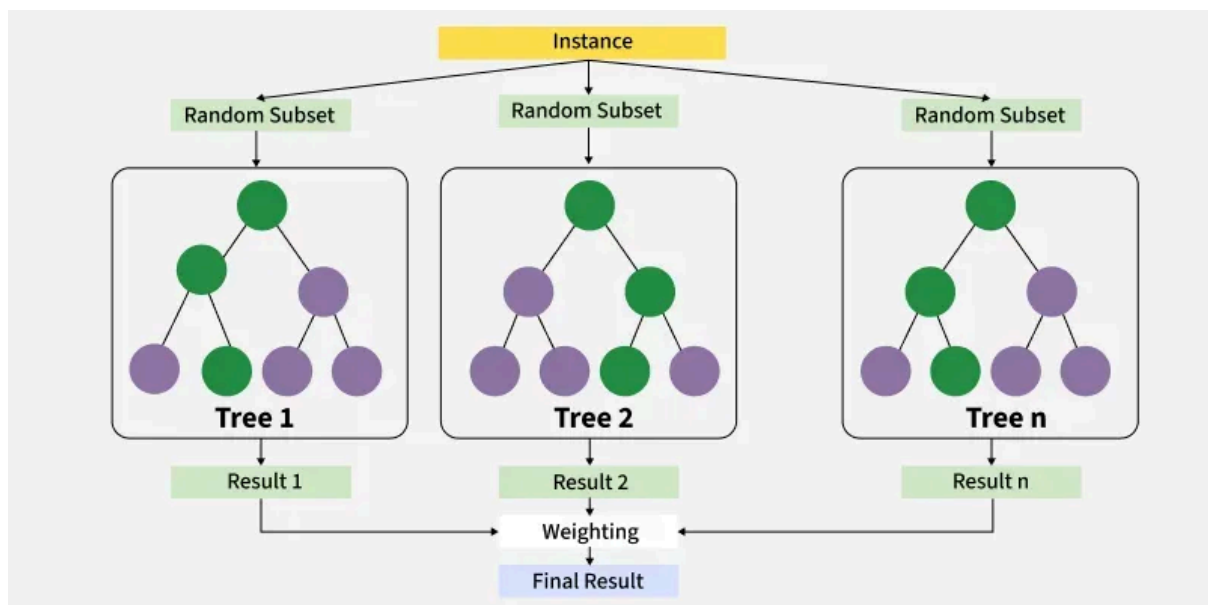
- Combines decision trees
  - Each new tree is trained to correct the previous one → **boosting**
- Parallel processing for high speed
- Allows for adjusting parameters to optimize performance based on the specific problem (would work well for us since we can change what it takes)

### **\*\* Gradient boost + Optimization techniques \*\***

- L1 and L2 Regularization
  - Prevents overfitting and increases generalization
- Model regularization
  - Built in feature for missing values
- Cross validation
- Parallelization
- Pruning
  - Depth first strategy
  - Deeper but more optimised trees
- Versatility
- Easy interface

## The algorithm

1. Base model
  - The first decision tree is trained on the data
  - Regression tasks → predicts the average of the target variable
2. Calculate the errors
  - Errors between the predicted and actual values are calculated
3. Train the next tree
  - The next tree is trained on the errors of the previous tree (corrects the previous errors)
4. Repeat the process
  - Process continues until a stopping criterion is met
5. Combine the predictions
  - The final prediction is the sum of the predictions from all the trees



## Advantages

- Good for large datasets
- Supports parallel processing and GPU acceleration → faster training
- Offers customizable parameters and regularization for fine-tuning
- Includes feature importance analysis for interpretability and explainability

## Disadvantages

- Computationally intensive
- Sensitive to noise or outliers, requiring careful data preprocessing
- **Prone to overfitting**, especially on small datasets or with too many trees.
- Model interpretability is limited compared to simpler methods
  - But does show feature importance

# Usage

## **When should we use it?**

- Classification problems, especially those related to real-world business problems.
- Problems in which the range or distribution of target values present in the training set can be expected to be similar to that of real-world testing data.
- Situations in which there are many categorical variables.
- large number of observations in training data.
- The number of features is smaller than the number of observations in training data.

## **When should we not use it?**

- The number of observations in training data is significantly smaller than the number of features.
- Computer vision
- Natural language processing
- Regression tasks that involve predicting a continuous output.
- Predicting increases in targets beyond the range present in the training data.
- Tasks involving extrapolation.

## **Installing links**

- <https://xgboost.readthedocs.io/en/stable/>
- [https://xgboost.readthedocs.io/en/stable/python/python\\_intro.html](https://xgboost.readthedocs.io/en/stable/python/python_intro.html)

## **References**

DataMListic. (2023, October 17). XGBoost Explained in Under 3 Minutes. YouTube.  
<https://www.youtube.com/watch?v=33fGfuleXw0>

GeeksforGeeks. (2021, September 18). XGBoost. GeeksforGeeks.  
<https://www.geeksforgeeks.org/xgboost/>

Nvidia. (2024). What is XGBoost? NVIDIA Data Science Glossary.  
<https://www.nvidia.com/en-us/glossary/xgboost/>

Adebisi, O. (2021). When to NOT use XGBoost? | Kaggle. Www.kaggle.com.  
<https://www.kaggle.com/discussions/general/196542>